

Speech recognition

D. Ligari 518592¹

¹ Machine Learning course, University of Pavia, Department of Computer Engineering (Data Science), Pavia, Italy

Github page: <https://github.com/DavideLigari01/speech-recognition.git>

Contact: davide.ligari01@universitadipavia.it

Date: May 11, 2023

Abstract—This report presents a lab activity focused on speech recognition. The task is to recognize the pronunciation of a single word from a list of 35 words, using a multilayer perceptron. The dataset used for the task is the Speech Commands Data Set, which includes 105,829 recordings of the 35 words, divided into training, validation, and test sets. Feature extraction has already been performed, and the features are spectrograms that have been made uniform in size. The lab activity encompasses various components, including the visualization of spectrograms, the application of feature normalization techniques, training a multilayer perceptron without hidden layers, and exploring different network architectures. To gain insights into the network's performance, a confusion matrix is constructed to summarize its behavior, and classification errors are thoroughly analyzed. The experiments are replicated using different feature normalization techniques, batch sizes and lambda values, in order to understand how these parameters affect the model's performance.

Keywords — MLP Neural network • Training • Speech recognition • Feature normalization • Confusion matrix • Spectrogram

1. MLP NEURAL NETWORK

The Multilayer Perceptron (MLP) neural network is a popular type of artificial neural network used in machine learning. It consists of interconnected layers of nodes or neurons that process data to produce predictions. MLPs employ activation functions, such as sigmoid, tanh, or ReLU, to introduce non-linearity and capture complex patterns. By adjusting the weights of these connections through a process called backpropagation, MLPs can learn from data and make accurate predictions. They are widely used in tasks like image recognition, speech processing, and natural language understanding.

2. VISUALIZE THE DATA

Gaining a deep understanding of the data being processed can greatly assist in designing an effective model. With this goal in mind, Figure 1 presents the spectrogram of the initial training sample, providing a visual representation of the data.




Fig. 1: Spectrogram of a sample of the dataset

3. BATCH SIZE SELECTION

Selecting an appropriate batch size is a critical factor in neural network training. Opting for a higher value facilitates smoother convergence towards the minimum point; however, each iteration becomes time-consuming as it needs to process a larger number of samples. Conversely, a very low batch size can lead to unstable convergence, requiring more iterations. On the positive side, each iteration is quicker since it processes only a limited number of samples.

The figure below illustrates the test accuracy for different epochs and batch sizes in a model without hidden layers. The model trained with a batch size of 10 exhibits gradual improvement but experiences oscillations. Differently, the 200-sized batch demonstrates more consistent progress, albeit at a slower pace compared to the 50 and 100-sized batches. Among these options, the 50-sized batch yields the highest accuracy, surpassing that of the 100-sized batch.




Fig. 2: Test accuracy for different batch sizes

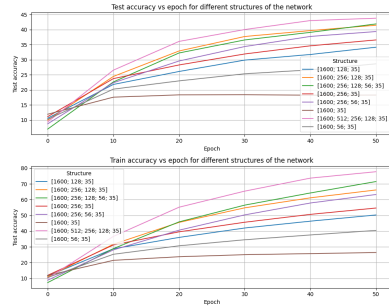


Fig. 3: Train and test accuracies for different network architectures

4. NETWORK ARCHITECTURE

a. Choice of the optimal lambda




Fig. 4: Test accuracy for different lambda values

b. Optimal network

Name server	ip
Google	8.8.4.4
Quad9	149.112.112.112
OpenDNS	208.67.220.220
Comodo Secure DNS	8.20.247.20

Table 1: Best neural network characteristics

5. MODEL ANALYSIS




Fig. 5: Confusion matrix of the best model

6. FEATURE NORMALIZATION

7. WEIGHTS VISUALIZATION

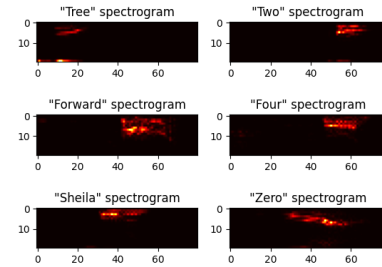


Fig. 6: Spectrogram of the most 3 misclassified words, on left the correct word, on right the confused one




Fig. 7: Spectrogram of a sample of the dataset after different normalizations




Fig. 8: Test accuracy for different normalizations

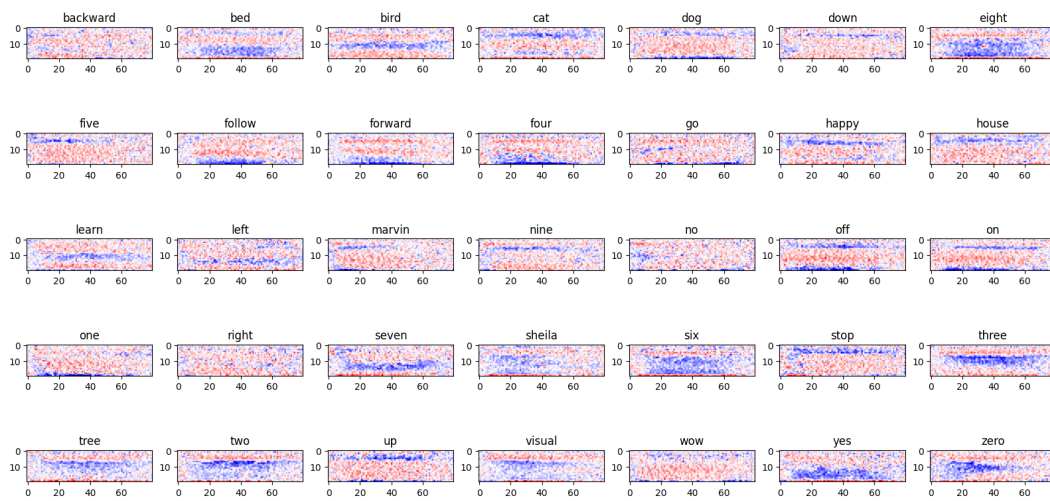


Fig. 9: spectrogram of the weights of the output layer